

Boxuan Shi

Pittsburgh, PA · eamesshi@gmail.com · [GitHub](#) · [LinkedIn](#)

EDUCATION

Carnegie Mellon University

Dec. 2026

Master of Information Systems Management — Business Intelligence and Data Analytics · GPA 3.88 / 4.0

- Coursework: Unstructured Data Analysis, A/B Testing, Introduction to Deep Learning, Causal Inference

Harbin Institute of Technology

Jul. 2025

Bachelor of Mgmt. in Digital Accounting · GPA 90.25 / 100 · National Scholarship (Top 0.2%)

- Coursework: Database, Econometrics, Operation Research, Probability Theory, Applied Statistics

PROFESSIONAL EXPERIENCE

Amgen — AI & Data, Remote, US

May 2026 – Aug. 2026

Data Scientist Intern, Applied AI Team

- Built a production-bound Travel & Entertainment budget forecasting model for the Finance organization across ~3,000 cost centers with 40 months of history; reduced forecast error ~8% on RMSE and ~17% on WAPE vs. baseline, and ~33% vs. human planners on both metrics.
- Engineered a threshold-based hybrid of two XGBoost models — standard L2 for high-spend transactions, and a log-L2 objective with sign classifier for the small, intermittent ones — to hold accuracy across a wide spend range; applied Tree SHAP to make each forecast explainable to Finance stakeholders.

Apple — Advertising, Shanghai, China

Feb. 2025 – Jul. 2025

Data Scientist Intern, Greater China Region Data Insight Team

- Diagnosed revenue leakage in an APAC Ads storefront by analyzing the revenue funnel and auction-level bid data with SQL (Athena, Snowflake) and PySpark; identified ~6% unfilled searches and surfaced high-potential keywords by scoring demand signals. Contributed to keyword coverage and brand protection strategies projected to drive ~15% uplift in annual revenue; reported to management across two milestones.
- Built a client-value model using linear regression with Scikit-learn on ~25 demographic and funnel features to estimate revenue headroom; visualized key insights with Matplotlib and Seaborn and delivered ranked client lists to Sales, improving pitch success rate by ~10%.

Deloitte — iBond, Shanghai, China

Jul. 2024 – Oct. 2024

Fintech Data Scientist Intern, Intelligent Stock Team

- Designed a Python/SQL ELTL pipeline processing 4,000+ stocks across 150 metrics daily for a market-monitor product; built a reusable multithreading decorator for concurrent query execution (5x speedup).
- Built credit scoring models in Python (NumPy, Pandas) for 18 industries using Genetic Algorithm optimization with custom constraints — category-level feature cardinality, weight caps, and a penalty for sign violations — to ensure interpretability; used Somers' D as the objective function, saving 10 person-days in project delivery.

PROJECTS

FIND — Internship Matching Platform

Nov. 2025 – Present

Independent Developer · <https://find-internship.org/>

- Developed an ML/AI-driven internship matching platform; gained 57 users within two weeks of beta. Built 66 crawlers monitoring ~19,000 active internship opportunities, with Prometheus and Grafana dashboards tracking system health, key metrics, and data drift.
- Integrated the DeepSeek LLM for resume parsing, job match reasoning, and cover letter generation — with prompt-injection protection and PII masking. Designed a bilingual semantic matching algorithm using a BGE-M3 encoder and pgvector, with a sigmoid to curve the matching score.
- Designed and ran a randomized landing-page experiment (50/50 block randomization, pre-specified ATE and HTE, covariate balance checks, power analysis) testing whether previewing sample output raises resume-upload intent; the effect was directionally positive with no moderation by privacy concern, and the sample-first onboarding flow shipped.

RESEARCH

Impact of Image in ESG Reports on ESG Performance

Nov. 2023 – Jun. 2025

ML Researcher · Supervised by Zhang, Qi at Harbin Institute of Technology, Weihai

- Built a multimodal analysis pipeline processing 11,000+ corporate ESG reports: automated PDF parsing for text and image extraction, and image-text semantic similarity computed with OpenAI CLIP.
- Discovered a U-shaped relationship between image density and ESG performance (inflection ~2.3 images/page), amplified by image-text alignment, using GLS regression with random effects in Stata.

SKILLS

Languages: Python, SQL, Java, Claude Code, Codex

Tools: Snowflake, PySpark, PostgreSQL, Scikit-learn, PyTorch, NumPy, Pandas, Matplotlib, Seaborn, Git

Analytics & Methodology: Causal Inference, Data Analysis, ELTL, Machine Learning, Natural Language Processing, Optimization